Lecture VI - Minimizing Split Orders

Applied Optimization with Julia

Dr. Tobias Vlćek
University of Hamburg

Introduction

E-Commerce Trends

Question: What are current trends in e-commerce?



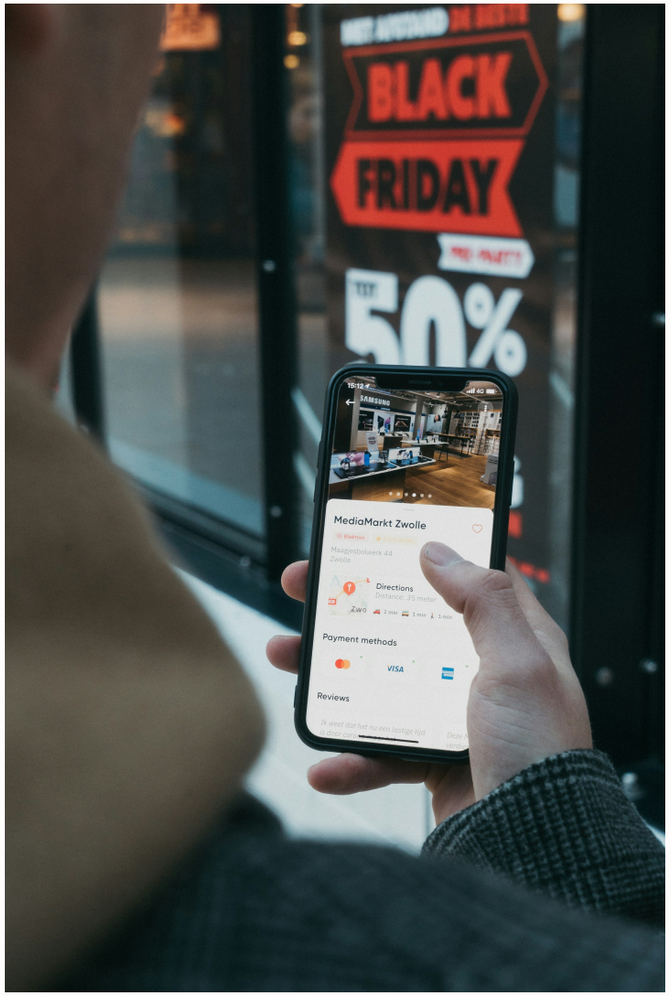
E-Commerce Sales

- E-Commerce sales **are growing fast**:
  - Products are no longer bound between borders
  - Product variety is rising
  - Consumer shopping patterns are shifting
  - Brick-and-mortar stores lose customers to the internet
  - Covid-19 accelerated this trend even more

Parcels Worldwide

- The number of parcels is rising:
  - 2014: 44 billion parcels (Pitney Bowes Inc. 2017)
  - 2019: 103 billion parcels (Pitney Bowes Inc. 2019)
  - 2026: forecast of 220 – 262 billion[1] (Pitney Bowes Inc. 2020)

Pressure on infrastructure



- Consumers nowadays expect **free, fast deliveries and returns**
- Existing warehouses have to store an increasing range of products
- Better customer service requires faster deliveries
- Incurred fulfillment costs depend on the number of parcels

Pressure on the environment



- Each parcel packaging consumes resources during production
- Every dispatched parcel to the customer **causes CO_2 emissions**
- In case of returns, more parcels cause more emissions

Learning Objectives

After this lecture, you will be able to:

- **Explain** split orders and why they matter in e-commerce
- **Build** a coappearance matrix from transactional data
- **Formulate and solve** the Quadratic Multiple Knapsack Problem (QMKP) in JuMP
- **Judge** when exact solvers reach their limits and heuristics are needed

Problem Structure

One Order, Two Parcels

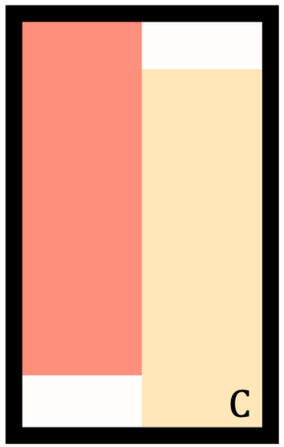
Question: Why did the shop send two?



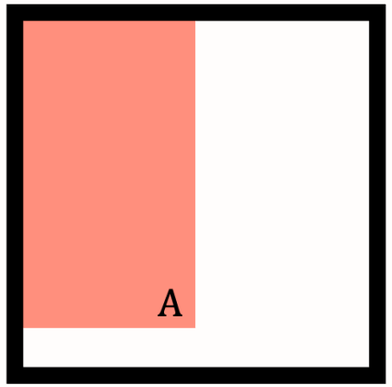
Photo: Thingsneverchange, Unsplash

Split Order

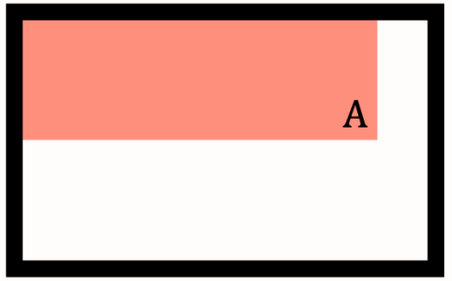
Customer's Order



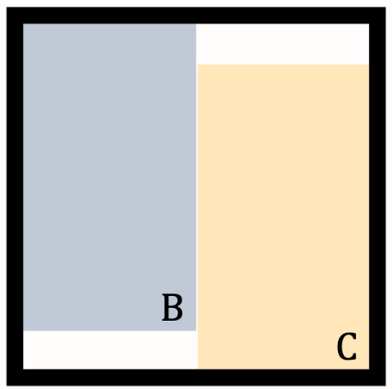
Warehouse ONE



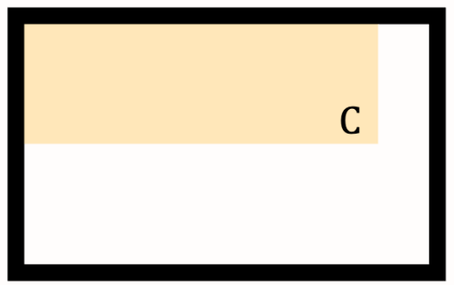
Parcel ONE



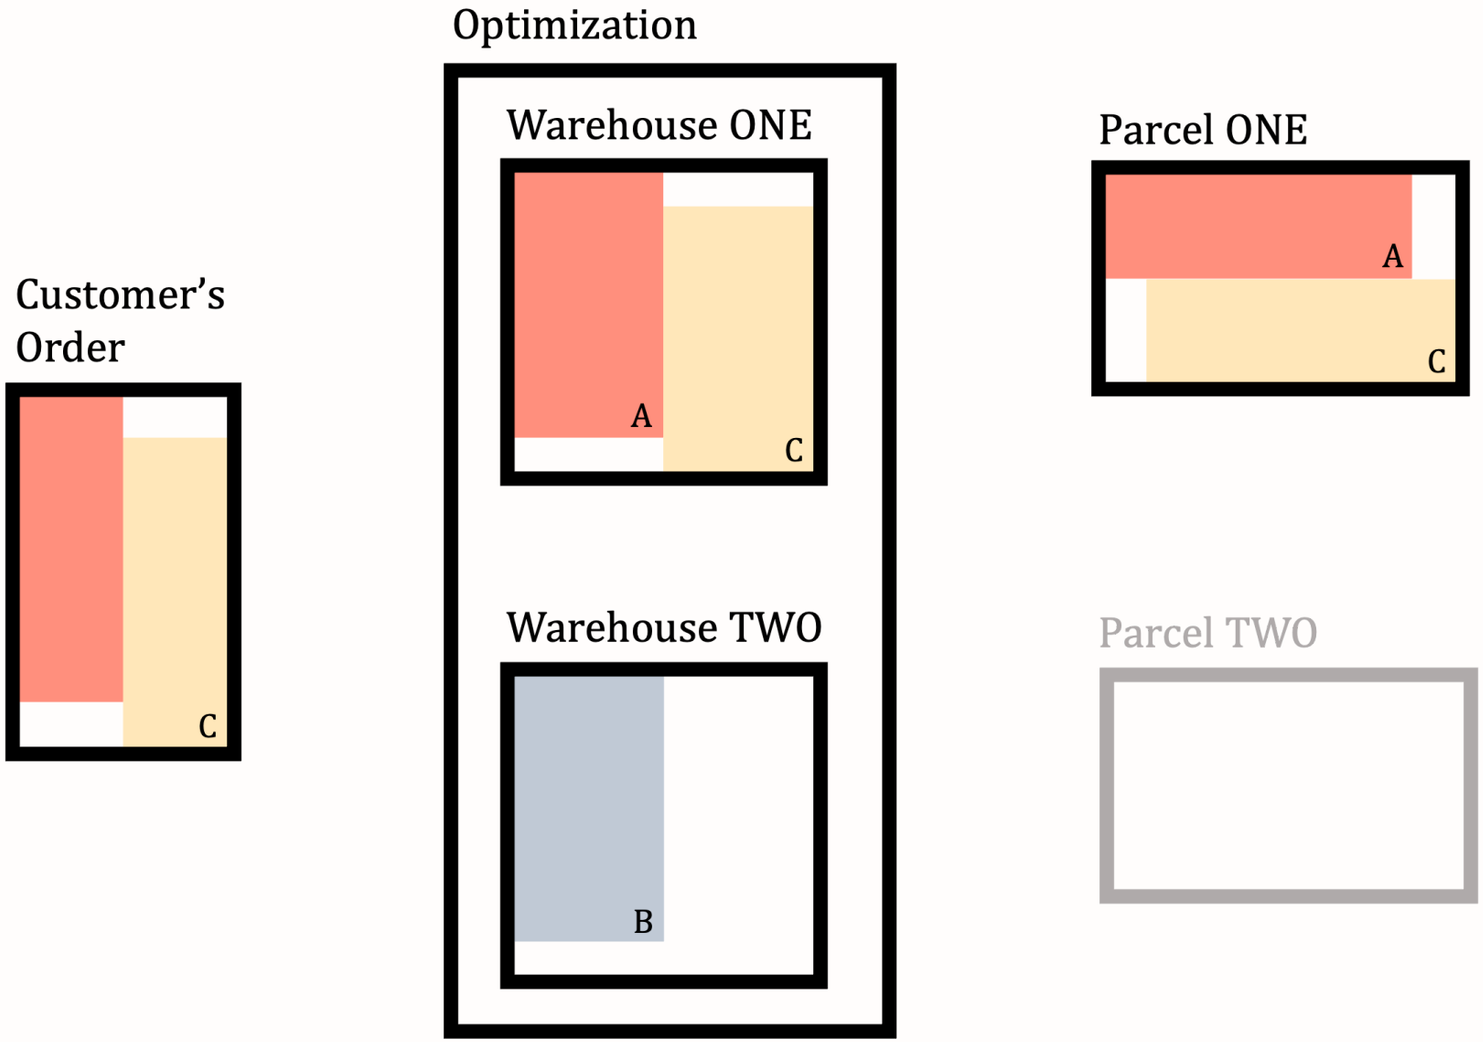
Warehouse TWO



Parcel TWO



No Split Order



Reason for Split Orders

Question: Why might they occur?

- Stock availability: Some products are **out of stock** at a warehouse and need to be fulfilled from another warehouse
- Capacity constraints: Some products are stored at **different warehouses** and need to be shipped from elsewhere

Impact of Split Orders

Question: What are the consequences?

- Higher shipping costs
- Increased packaging material
- More CO_2 emissions
- **Higher operational complexity**
- Lower customer satisfaction

Mitigations?

Question: What are possible mitigations?

- Consolidation: Ship to a central warehouse before dispatch
- Cross-docking: Transfer goods from inbound to outbound transport with little or no storage in between
- Transshipment: Ship between warehouses before delivery
- Co-allocation: Predict **coappearance of products** and allocate them to the **same warehouse**

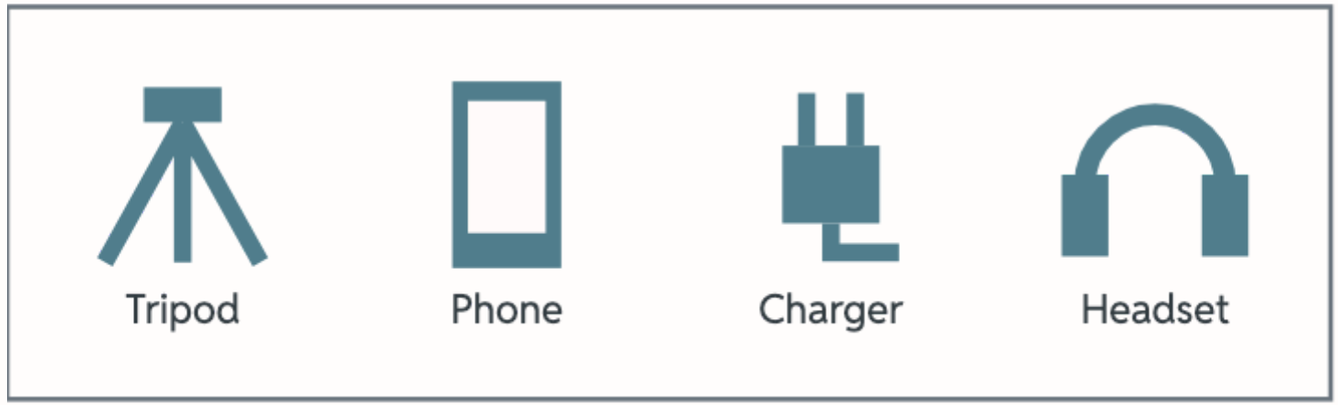
Your Turn

An online shop stores 8 products in Hamburg and Berlin, 4 places each; an order with products from both ships as two parcels. Drag a product onto one in the other warehouse to swap them, and split as few parcels as possible.

Hamburg



Berlin



Bought together



You **18 split parcels**

Ready: press Optimize to compare with the best plan.

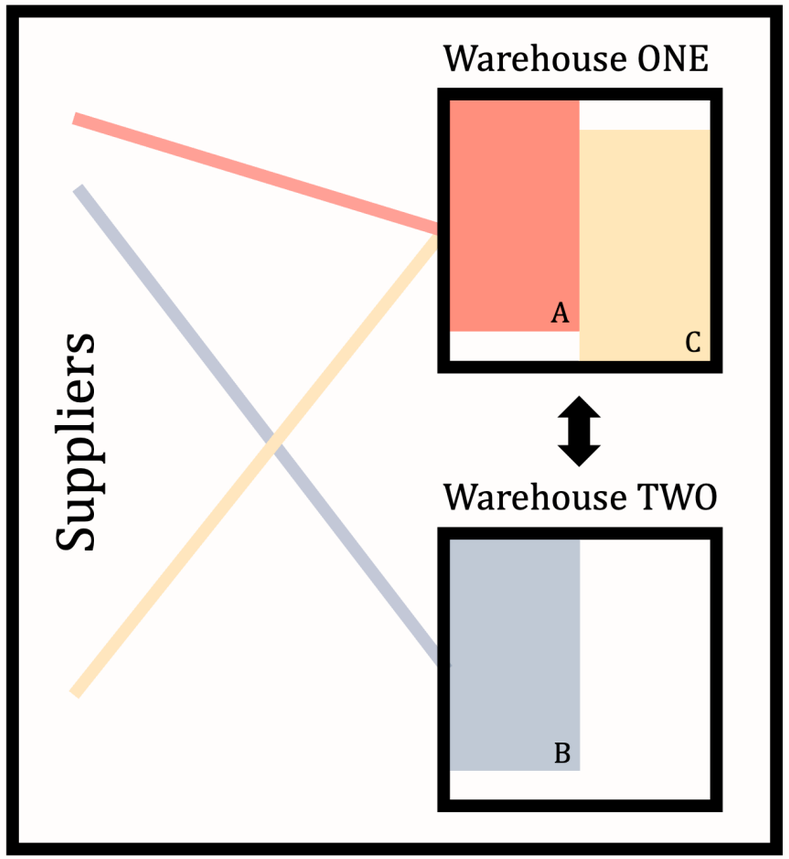
Optimize

Reset

New

Structure - Version 1

Optimizing Co-allocation



Question: What could be our objective?

We aim to improve the SKU[1]-warehouse allocation to minimize the number of split parcels resulting from SKUs being stored in different warehouses.

Available Sets

Question: What could be the sets here?

- $\mathcal{I}$ - Set of products indexed by $i \in \{1, 2, \ldots, |\mathcal{I}|\}$
- $\mathcal{K}$ - Set of warehouses indexed by $k \in \{1, \ldots, |\mathcal{K}|\}$
- $\mathcal{M}$ - Set of customer orders $m \in \{1, 2, \ldots, |\mathcal{M}|\}$

Available Parameters

Question: What are possible parameters?

- c_k - Storage space of warehouse $k \in \{1, \ldots, |\mathcal{K}|\}$, measured in number of SKUs
- $\boldsymbol{T} = (t_{m,i})$ - Past customer orders for SKUs

Question: What could the transactional data look like?

Transactional Data

Example of

	A	B	C	D
1	1	1	1	0
2	1	1	1	0
3	1	1	0	0
4	1	0	0	1
5	1	0	0	1
6	1	0	0	1
7	1	0	0	1
8	0	0	1	1

Past vs. Future

- The **transactional data** is based on **past orders**
- It is a **binary matrix** of customer orders and SKUs
- We use this data to **assume** future co-occurrence
  - **Past co-occurrence predicts future co-occurrence**

Question: What is your opinion on the assumption?

Split-Order Minimization

Question: What could be our decision variable/s?

ⓘ We have the following sets:

- - Set of products indexed by
- - Set of warehouses indexed by
- - Set of customer orders

- - 1, if is stored in , 0 otherwise
- - 1, if SKU is shipped from warehouse for customer order , 0 otherwise

Integer Programming Model

- Catalán and Fisher (2012) created an **integer model**
- Number of SKUs of E-Commerce retailers can easily be **between 10,000 - 100,000**
- The number of customer orders needed for “stable” results is one order of magnitude higher: **100,000 - 10,000,000**

Question: Does anyone have an idea what this could mean?

Why is Version 1 so hard?

- Objective: **minimize the number of warehouses** that ship parcels for each customer order
- Requires : one variable per order, SKU **and** warehouse
- binary variables

> ⚠ **Variables explode with the instance size**
>
> 10,000 SKUs, 1,000,000 orders and 2 warehouses already lead to binary variables !

Implementation Challenges

- Small instance with 10 SKUs and 1000 customer orders
- CPLEX 20.1.0 needs 3100 seconds to solve the problem
- Computation time scales exponentially
- Not applicable in real world applications!

Any idea what could be done?

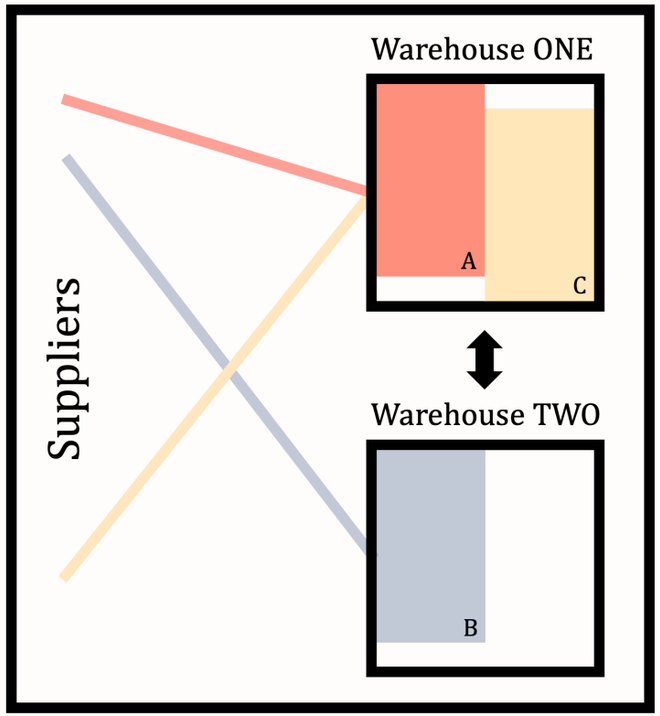
Structure - Version 2

Heuristic Approach

- **Heuristic**: Fast, but not necessarily optimal
- **Approximation**: Not guaranteed to be optimal, but close
- **Computational Effort**: Reasonable even for large instances

> **Different view on the problem**
>
> Focus on the warehouses and the coappearance of SKUs! Discard the exact information about the customer orders.

Objective



Question: What could be the objective?

Maximize the coappearance of products that are often part of the same customer orders.

Transaction Matrix

```
T = [
    1 1 1 0;
    1 1 1 0;
    1 1 0 0;
    1 0 0 1;
    1 0 0 1;
    1 0 0 1;
    1 0 0 1;
    0 0 1 1
]

# Create the coappearance matrix
Q = T' * T
println("Coappearance matrix Q:")
display(Q)
```

```
Coappearance matrix Q:
4×4 Matrix{Int64}:
 7  3  2  4
 3  3  2  0
 2  2  3  1
 4  0  1  5
```

Coappearance Matrix

- is a **symmetric matrix**
- Proposed by Catalán and Fisher (2012)
- where
- shows how often and appear **in the same order**

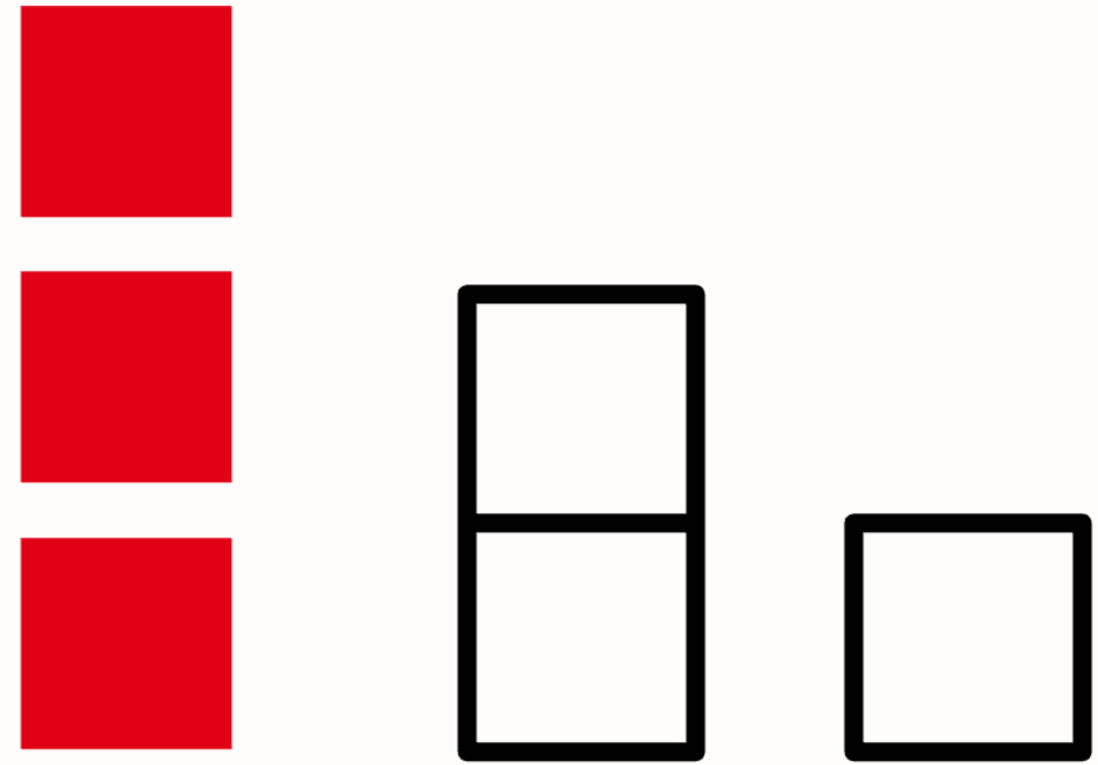
Question: What do the principal diagonal values tell us?

- In how many orders each SKU appeared
- As is binary, counts orders, **not units sold**

How to approach the problem?

- Greedy Heuristic[1]: Allocation based on matrix
- Mathematical Model[2]: Maximizes coappearance
- GRASP[3]: Good on small instances
- **New**: Max. coappearance with non-linear solver
- **New**: Heuristic based on Chi-Square Tests

Basic Setting



SKU i Warehouse k

Available Data (Version 2)

Question: What could be the sets?

- - Set of products indexed by
- - Set of warehouses indexed by

> ⚠ **No customer order information is needed!**
>
> We can focus on the SKUs and the warehouses, making the problem **much smaller**!

Available Parameters

Question: What are possible parameters?

- - Storage space of warehouse , measured in number of SKUs
- - Coappearance matrix

> ⚠ **Transactional Data replaced**
>
> Instead of the transactional data, we just **use the coappearance matrix** in our model!

Model Formulation

Decision Variables?

ⓘ We have the following sets:

- - Set of products indexed by
- - Set of warehouses indexed by

ⓘ Our objective is to:

Maximize the coappearance of products that are often part of the same customer orders.
In more mathematical terms: Maximize the sum of all unique pair-wise values of all SKUs stored in the same warehouse.

Question: What could be our decision variable/s?

Decision Variables

- - 1, if SKU is stored in , 0 otherwise

> **Only one variable per SKU and warehouse!**
>
> As we don’t need the customer order information, we only need to make a decision for each SKU and warehouse pair!

Decision Variable in Julia

Question: How could we formulate the variable in Julia?

```julia
using JuMP, SCIP # SCIP is a non-commercial solver for quadratic models

warehouses = ["Hamburg", "Berlin"] # Add warehouses as a vector
skus = ["A", "B", "C", "D"] # The four SKUs from our transactional data

warehouse_model = Model(SCIP.Optimizer)
```

```julia
@variable(warehouse_model, X[i in skus, k in warehouses], Bin)
```

```
2-dimensional DenseAxisArray{JuMP.VariableRef,2,...} with index sets:
    Dimension 1, ["A", "B", "C", "D"]
    Dimension 2, ["Hamburg", "Berlin"]
And data, a 4×2 Matrix{JuMP.VariableRef}:
 X[A,Hamburg]  X[A,Berlin]
 X[B,Hamburg]  X[B,Berlin]
 X[C,Hamburg]  X[C,Berlin]
 X[D,Hamburg]  X[D,Berlin]
```

Objective Function

ⓘ **We need the following:**

- - 1, if SKU is stored in , 0 otherwise
- - Coappearance of SKU and

ⓘ **Our objective is to:**

Maximize the sum of all unique pair-wise values of all SKUs stored in the same warehouse. Note, that this is a **quadratic objective function**!

Question: What could the objective function look like?

Quadratic Objective Function

This is a quadratic objective function!

The quadratic terms are . This objective function is based on the **Quadratic Multiple Knapsack Problem (QMKP)**, formulated by Hiley and Julstrom (2006).

Objective Function in Julia

Question: How could we formulate this in Julia?

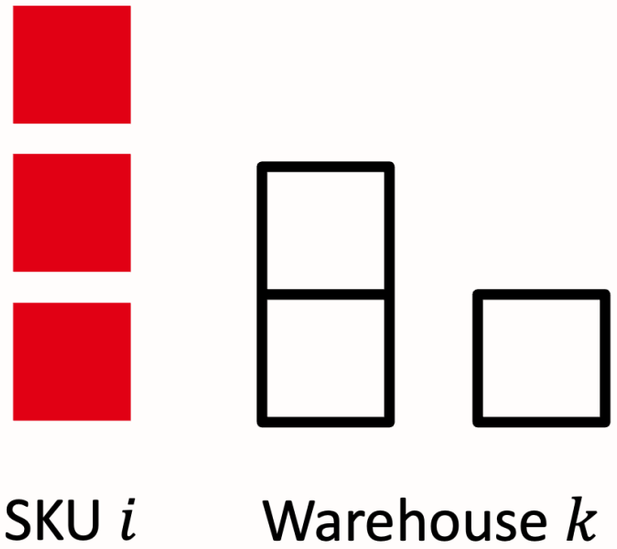
```julia
# We reuse the matrix Q = T' * T computed from our transactional data

@objective(warehouse_model,
    Max,
    sum(
        X[skus[i], warehouses[k]] * X[skus[j], warehouses[k]] * Q[i,j]
        for i in 2:length(skus)
        for j in 1:i-1
        for k in 1:length(warehouses)
    )
)
```

```
3 X[B,Hamburg]*X[A,Hamburg] + 3 X[B,Berlin]*X[A,Berlin] + 2 X[C,Hamburg]*X[A,Hamburg] + 2
X[C,Berlin]*X[A,Berlin] + 2 X[C,Hamburg]*X[B,Hamburg] + 2 X[C,Berlin]*X[B,Berlin] + 4
X[D,Hamburg]*X[A,Hamburg] + 4 X[D,Berlin]*X[A,Berlin] + X[D,Hamburg]*X[C,Hamburg] +
X[D,Berlin]*X[C,Berlin]
```

Constraints

What constraints?



Question: What constraints?

- Allocate each SKU at least once
- Warehouses have a finite capacity
- Capacity is not exceeded

Single Allocation Constraint?

The goal of this constraint is to:

Ensure that each SKU is allocated at least once.

We need the following variable:

- - 1, if SKU is stored in , 0 otherwise

Question: What could the constraint look like?

Single Allocation Constraint

> ⓘ Remember, this is the variable:
>
> - - 1, if SKU is stored in , 0 otherwise

Question: How could we change the constraint to ensure that each SKU is allocated only once?

Question: How could we add the constraint in Julia?

Single Allocation in Julia

```julia
@constraint(warehouse_model, single_allocation[i in skus],
    sum(X[i, k] for k in warehouses) >= 1
)
```

```
1-dimensional DenseAxisArray{JuMP.ConstraintRef{JuMP.Model,
MathOptInterface.ConstraintIndex{MathOptInterface.ScalarAffineFunction{Float64},
MathOptInterface.GreaterThan{Float64}}, JuMP.ScalarShape},1,...} with index sets:
    Dimension 1, ["A", "B", "C", "D"]
And data, a 4-element Vector{JuMP.ConstraintRef{JuMP.Model,
MathOptInterface.ConstraintIndex{MathOptInterface.ScalarAffineFunction{Float64},
MathOptInterface.GreaterThan{Float64}}, JuMP.ScalarShape}}:
 single_allocation[A] : X[A,Hamburg] + X[A,Berlin] ≥ 1
 single_allocation[B] : X[B,Hamburg] + X[B,Berlin] ≥ 1
 single_allocation[C] : X[C,Hamburg] + X[C,Berlin] ≥ 1
 single_allocation[D] : X[D,Hamburg] + X[D,Berlin] ≥ 1
```

Capacity Constraints?

The goal of these constraints is to:

Ensure that the capacity of each warehouse is not exceeded.

We need the following variables and parameters:

- - 1, if SKU is stored in , 0 otherwise
- - Storage space of warehouse

Question: What could the second constraint be?

Capacity Constraints

And that's basically it!

Question: How could we add the second constraint in Julia?

Capacity Constraints in Julia

```julia
capacities = Dict("Hamburg" => 2, "Berlin" => 2) # Add capacities

@constraint(warehouse_model, capacity[k in warehouses],
    sum(X[i, k] for i in skus) <= capacities[k]
)
```

```
1-dimensional DenseAxisArray{JuMP.ConstraintRef{JuMP.Model,
MathOptInterface.ConstraintIndex{MathOptInterface.ScalarAffineFunction{Float64},
MathOptInterface.LessThan{Float64}}, JuMP.ScalarShape},1,...} with index sets:
    Dimension 1, ["Hamburg", "Berlin"]
And data, a 2-element Vector{JuMP.ConstraintRef{JuMP.Model,
MathOptInterface.ConstraintIndex{MathOptInterface.ScalarAffineFunction{Float64},
MathOptInterface.LessThan{Float64}}, JuMP.ScalarShape}}:
 capacity[Hamburg] : X[A,Hamburg] + X[B,Hamburg] + X[C,Hamburg] + X[D,Hamburg] ≤ 2
 capacity[Berlin] : X[A,Berlin] + X[B,Berlin] + X[C,Berlin] + X[D,Berlin] ≤ 2
```

QMK Model

subject to:

QMK Model in Julia

```julia
set_attribute(warehouse_model, "display/verblevel", 0) # Hide solver output
optimize!(warehouse_model)

println("The optimal objective value is: ", objective_value(warehouse_model))
println("The optimal solution is: ", value.(X))
```

```
The optimal objective value is: 6.0
The optimal solution is: 2-dimensional DenseAxisArray{Float64,2,...} with index sets:
    Dimension 1, ["A", "B", "C", "D"]
    Dimension 2, ["Hamburg", "Berlin"]
And data, a 4×2 Matrix{Float64}:
 1.0  0.0
 0.0  1.0
 0.0  1.0
 1.0  0.0
```

Question: What does this value tell us?

Interpreting the Solution

- Total capacity () **equals** the number of SKUs
- Hence, each SKU is stored **exactly once**
- Optimal: **A and D** share a warehouse, so do **B and C**

> ⓘ **The objective value is 6**
>
> It sums the coappearances of SKUs sharing a warehouse: . It does **not** directly count the avoided split orders!

Model Characteristics

Characteristics

- Is the model formulation **linear/ non-linear?**
- What kind of **variable domain** do we have?
- Do we know the **split-orders** based on the **objective value?**
- Why **couldn’t we use HiGHS** as solver?

Choosing a solver

- Identify **problem structure**, e.g. LP, MIP, NLP, MIQP, MIQCP, ...
- What is the **size** of the problem?
- Is a **commercial** solver needed?

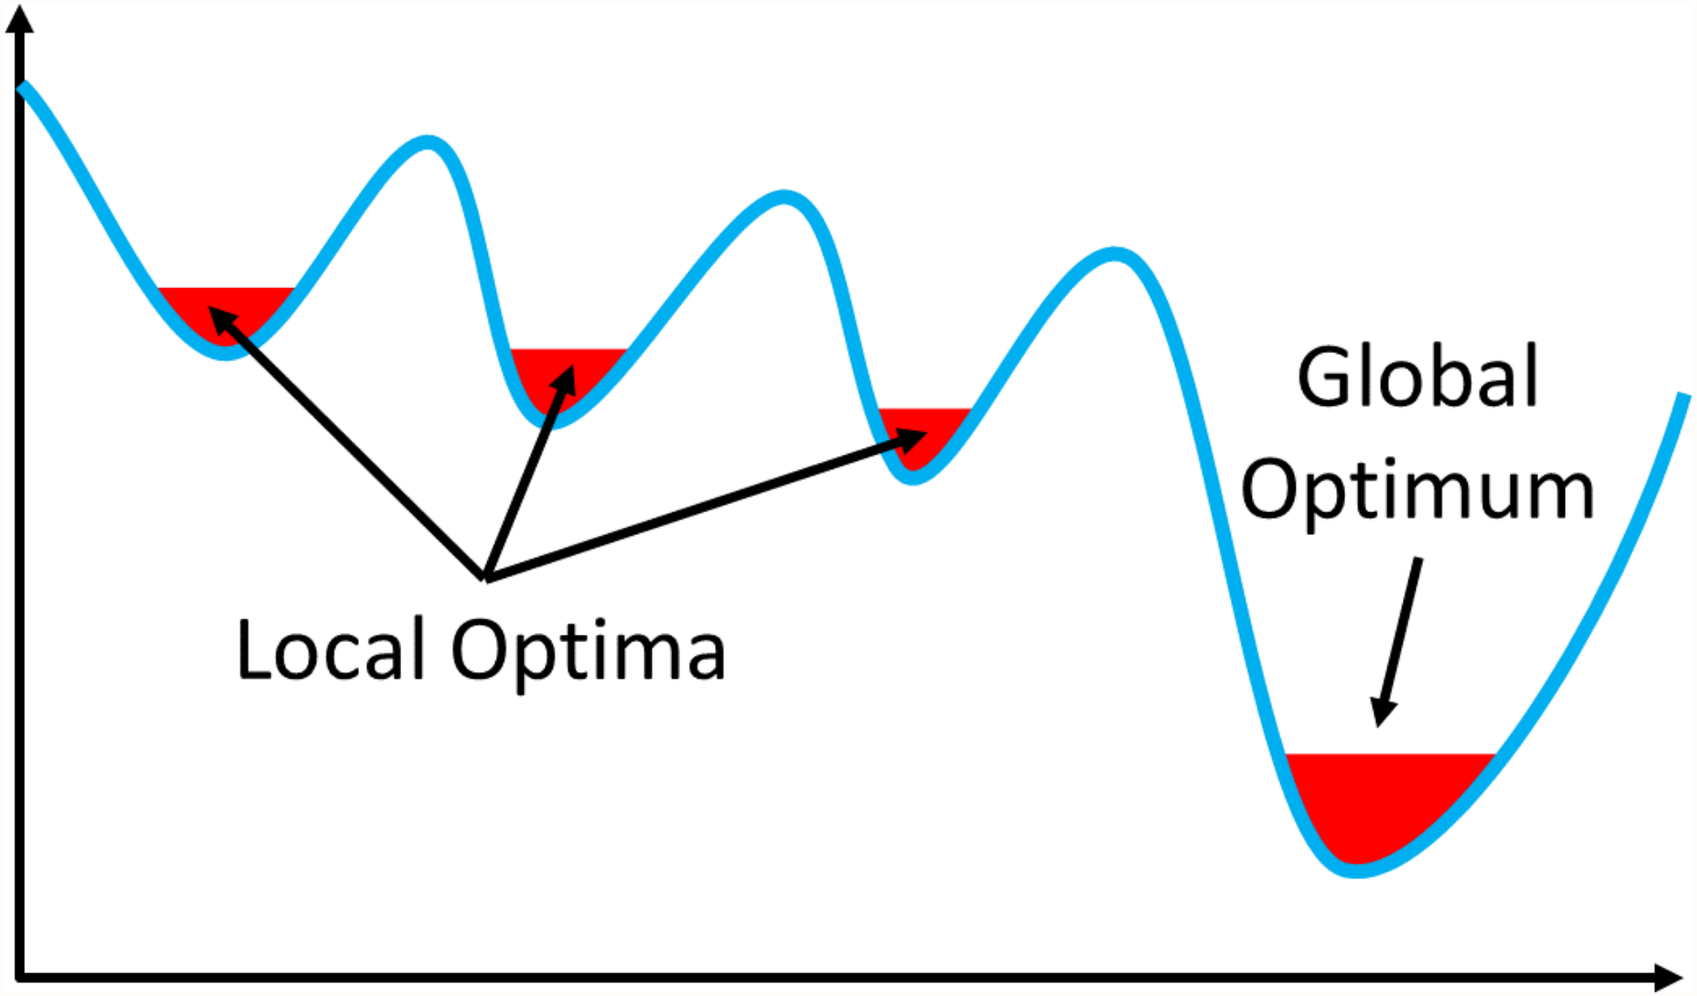
> ⓘ **Commercial Solvers**
>
> Commercial solvers are **faster** and **more robust** than open source solvers but also **more expensive**. During your studies, you can use most of them for free though! Nonetheless, we will only use open source solvers in this course.

Our Problem Class

- Objective: quadratic terms
- Constraints: all linear
- Variables: all binary

> ⓘ **Mixed-Integer Quadratic Program (MIQP)**
>
> Our model is a MIQP, as only the objective is quadratic. It is not a MIQCP, which would also allow quadratic constraints. MIQCP solvers, e.g. SCIP, subsume MIQPs and can thus solve our model.

Global vs Local Optimality



Solver Comparison[1]

SKUs	SCIP	HiGHS	Gurobi
100	118s	X	~400s
1,000	1,011s (18%)	X	~200s (2%)
10,000	X	X	X

Solver Comparison: Takeaways

- **SCIP**: Best open-source for MIQP/MIQCP, limited scalability
- **Commercial solvers**: Better but still fail on 10,000+ SKUs
- **Conclusion**: Heuristics necessary for realistic problems!

Model Assumptions

Questions: On model assumptions

- What assumptions have we made?
- Problem with allocating SKUs to multiple warehouses?
- What else might pose a problem in the real world?

> ⓘ **Duplicates can inflate the objective**
>
> With slack capacity, the model may store one SKU in several warehouses. A shared pair then counts once per warehouse, raising the objective without avoiding any further splits.

Model Limitations

- Stock-outs: Model assumes **perfect availability**
  - Real stockouts create splits despite optimal allocation
- Storage costs: No differentiation between costs
- Demand dynamics: Weather, promotions, seasonality
- SKU sizes: Uniform capacity assumption unrealistic
- Shipping costs: No consideration of distance or weight

Practical Limitations

Allocation changes:

- Frequent reallocation **disrupts operations**
- Physical inventory movements costly
- Need for change minimization

Workload imbalance:

- Optimal allocation may **overload** some warehouses
- High-frequency SKUs cluster together
- Trade-off: splits vs. balance

Impact

Can this be applied?

Problem Size is Crucial

- Up to 1,000 SKUs → commercial solvers
- More than 1,000 SKUs → heuristics
- For example, the **CHI** heuristic

> **CHI-Heuristic**
>
> Detect dependencies between products and allocate them accordingly, as products within orders can have dependencies and products are bought with different frequencies!

CHI Heuristic

Statistical Independence

Question: When are two SKUs independent?

If purchasing one doesn't affect purchasing the other:

Example: If appears in 10% of orders and in 20%:

- Independent: Appear together in of orders
- Dependent: Appear together in **5% of orders** → customers buy them together!

Chi-Square Test Intuition

Use chi-square tests to detect if SKUs are dependent:

- Calculate **expected coappearances** under independence
- Compare with actual coappearances
- If difference is significant → SKUs are dependent
- Allocate dependent SKUs to **same warehouse**!

From Tests to Allocations

Two-Phase Approach

Phase 1: Build allocation using dependencies
Phase 2: Refine with local search

Result: Works on large instances, 50,000 SKUs in 10 minutes!

In Practice

Real-World Challenges

Question: What challenges might arise in practice?

- New products: Allocate SKUs without historical data?
- Dynamic inventory: SKUs appear and disappear
- Different sizes: SKUs consume varying storage space
- Computational limits: Problems with 50,000+ SKUs

Case: A Large Online Retailer

Question: Which of 100,000+ SKUs goes to which warehouse?



Photo: Alberto Rodríguez, Unsplash

Rolling Horizon Approach

In the following case study (Vlćek and Voigt 2024), we:

- used a **5-week training window** for allocation decisions
- updated allocations **weekly** based on recent patterns
- balanced **stability** vs. **responsiveness**



In the case study, SKUs were grouped by **category-brand combinations** to reduce the problem size while maintaining allocation quality. New SKUs inherit the allocation of their cluster!

Case Study

Problem characteristics:

- 100,000+ SKUs
- Several million orders
- Multiple warehouses
- Storage capacity constraints

Heuristics compared[1]:

- **CHI**: Chi-Square tests (Vlćek and Voigt 2024)
- **GP, GO, GS, BS**: Greedy (Catalán and Fisher 2012)
- **RA**: Random allocation

Implementation Results

Retailer's status quo: 6.95% of all orders were split

Theoretical reduction of these splits:

- CHI: **82.25%** reduction
- BS: **62.63%** reduction
- GS: **59.16%** reduction

Conclusion

- Splits are of no benefit, except faster customer deliveries
- **Increase workload, packaging and shipping costs**
- Mathematical Optimization of “full” problem not solvable
- CHI Heuristic close to mathematical optimization

> ⓘ And that’s it for today’s lecture!
>
> We now have covered the Quadratic Multiple Knapsack Problem and are ready to start solving some tasks in the upcoming tutorial.

Questions?

Literature

Literature I

References for this lecture:

- Catalán and Fisher (2012): Integer model and greedy heuristics
- Hiley and Julstrom (2006): The Quadratic Multiple Knapsack Problem
- Zhu et al. (2021): Mathematical model and GRASP
- Vlćek and Voigt (2024): CHI heuristic and case study

Literature II

For more interesting literature to learn more about Julia, take a look at the literature list of this course.

References

Catalán, Andrés, and Marshall Fisher. 2012. “Assortment Allocation to Distribution Centers to Minimize Split Customer Orders.” *SSRN Electronic Journal*, ahead of print. https://doi.org/10.2139/ssrn.2166687.

Hiley, Amanda, and Bryant A. Julstrom. 2006. “The Quadratic Multiple Knapsack Problem and Three Heuristic Approaches to It.” In *Proceedings of the 8th Annual Conference on Genetic and Evolutionary Computation*, edited by M. Keijzer. Association for Computing Machinery. https://doi.org/10.1145/1143997.1144096.

Pitney Bowes Inc. 2017. “Pitney Bowes Parcel Shipping Index Reveals 48 Percent Growth in Parcel Volume since 2014.” https://www.businesswire.com/news/home/20170830005628/en/Pitney-Bowes-Parcel-Shipping-Index-Reveals-48.

Pitney Bowes Inc. 2019. “Pitney Bowes Parcel Shipping Index Reports Continued Growth Bolstered by China and Emerging Markets.” https://www.businesswire.com/news/home/20191010005148/en/.

References (2/2)

Pitney Bowes Inc. 2020. “Pitney Bowes Parcel Shipping Index Reports Continued Growth as Global Parcel Volume Exceeds 100 billion for First Time Ever.” https://www.businesswire.com/news/home/20201012005150/en/.

Vlćek, Tobias, and Guido Voigt. 2024. *Optimizing SKU-Warehouse Allocations to Minimize Split Parcels in E-Commerce Environments*. Working paper, to be submitted to Manufacturing & Service Operations Management.

Zhu, Shan, Xiangpei Hu, Kai Huang, and Yufei Yuan. 2021. “Optimization of Product Category Allocation in Multiple Warehouses to Minimize Splitting of Online Supermarket Customer Orders.” *European Journal of Operational Research* 290 (2): 556–71. https://doi.org/10.1016/j.ejor.2020.08.024.